David Schwartzman

10/30/2025

GAN Model Discussion

1) **Trade-off between Image Quality and Diversity**:

In GANs, there's a balance between generating high-quality images and ensuring diversity in the outputs. In my case, the images are noisy and unclear, likely because the model hasn't had enough training to capture the details required for realism. When GANs are under-trained, they often struggle with both quality and diversity, meaning that while the images may vary, they lack the clear, recognizable features of bedrooms.

Since the images are not yet recognizable, it's hard to discuss the diversity accurately. But with more training time, I might see an improvement where the GAN captures different bedroom styles more distinctly. The trade-off here may lean heavily towards generating clearer images over diversity until the model reaches a baseline quality.

2) **Potential Enhancements**:

Adjusting Hyperparameters: Experimenting with hyperparameters like learning rate, batch size, or the architecture's depth might improve the output. For instance, a smaller batch size could allow the model to learn finer details, while more epochs could help improve the image clarity.

Regularization Techniques: Adding techniques like label smoothing or spectral normalization may stabilize training and reduce noise, helping the GAN to produce images with better-defined features.

3) **Practical Applications**:

Generative models like GANs are valuable for applications such as interior design, where AI-generated room designs can help inspire layouts and decor ideas. GANs can also be used in video game design, virtual reality, and movie production, where diverse room environments are needed for immersive experiences.

Beyond entertainment, GANs can also assist in fields like real estate and architecture, where AI-generated rooms can visualize home staging or preview design choices before actual implementation.

4) **Extra Credit Research Problem**:

To quantitatively evaluate my images, I could consider using the Inception Score (IS) or Frechet Inception Distance (FID). Both metrics provide a numerical assessment of GAN outputs by measuring the similarity of generated images to real ones:

Inception Score (IS): This score evaluates image quality and diversity by checking how confidently a classifier model identifies generated images, with higher confidence indicating better quality.

Frechet Inception Distance (FID): FID compares the distributions of real and generated images, assessing both quality and similarity to the real image dataset. Lower FID values suggest higher quality and diversity, as they indicate generated images are closer in distribution to the original dataset.